Acta Crystallographica Section A
Foundations of Crystallography
ISSN 0108-7673

The Selling closure: a complete operator set for reindexing, and the price of pseudo-symmetry

Petrus H. Zwart *

Molecular Biophysics and Integrated Bioimaging Division, Lawrence Berkeley National Laboratory, 1 Cyclotron Road, Berkeley, CA 94720, USA. Correspondence e-mail: phzwart@lbl.gov

A lattice with a right angle between vectors of its reduced superbase has several Selling-reduced cells, and two measurements of one crystal can reduce to different ones. The integer matrices from a reduced cell to itself form the lattice point group, and those between two reduced cells form one coset of it. Every such matrix, for any lattice, lies in one fixed set of $6096 = 48 \times 127$ matrices with entries up to ± 3, and which of them occur depends only on which conorms vanish. The proof is a short lemma and an exhaustive enumeration over the 38 boundaries of the cone of reduced conorms, checked in Lean, the finite steps by compiled evaluation. For a measured cell, the superbases that are reduced for some metric within a strain tolerance are finitely many and are enumerated. Reindexing a new cell against such a reference returns exactly the proper operators whose price lies within a chosen band, the price being the affine-invariant distance between the two metric tensors after the best common scale; the true operator is among them whenever its price is. A certificate says when the fixed set of 6096 matrices alone suffices. The candidates lie between two cosets of the reference's near-symmetries, these are bounded by the symmetries of the nearest symmetric lattice, and a pseudo-twofold costs exactly twice its distance to the metrics it would leave invariant, about 5% per degree of Le Page obliquity. For pairs of deposited cells of one crystal form in the PDB, taken as adjacent entries of one space group, the cell leaves a priced alternative indexing open for about one pair in thirteen at a band of 7.4%, and ranks a wrong indexing first for 0.1 to 1.2% of pairs, depending on how relabellings of axes are counted. These measure what the cell metric cannot decide, not how often data were misindexed.

© 0000 International Union of Crystallography
Printed in Singapore – all rights reserved

1. Introduction

Take a primitive monoclinic lattice with $a = 30$, $b = 40$, $c = 50$ Å and $\beta = 100°$. Its Selling-reduced cell has two right angles between vectors of the reduced superbase, and a measured cell lies slightly to one side or the other of each. One data set therefore reduces to the cell 30, 40, 50 Å, $\beta = 100°$, and another, equally correctly, to 30, 40, 53.7 Å, $\beta = 113.4°$, built on a and $a + c$ (Section 3.2). The matrix between the two is not a symmetry of the lattice. A program that reports the Niggli cell (Niggli, 1928; Santoro & Mighell, 1970; Křivý & Gruber, 1976) meets the same multiplicity in another form: the cell on a and $a + c$ is never a Niggli cell, but at each right angle the Niggli cell changes between its two types, with three acute or no acute angles (Gruber, 1973; Andrews *et al.*, 1980), and two measurements of one crystal can be reported in Niggli cells related by a matrix that is not a symmetry either. Data-processing programs (Kabsch, 2010; Winter *et al.*, 2018; Liebschner *et al.*, 2019) find such matrices by trying a list: small integer matrices between the two reduced cells, the matrices that unit-cell similarity searches generate (Grosse-Kunstleve *et al.*, 2002), or the operations of the lattice group found by a symmetry search (Le Page, 1982; Sauter *et al.*, 2004). On ordinary cases these searches succeed. This paper asks which integer matrices can relate two reduced cells, how to enumerate every candidate within a stated tolerance, and how to price the pseudo-symmetric alternatives that a search returns.

The multiplicity has a classical description, and this paper uses its established language. A superbase is four lattice vectors summing to zero; it is obtuse, or Selling-reduced, when all six pairwise dot products are non-positive (Selling, 1874; Conway & Sloane, 1992), and the obtuse superbase is unique except when a dot product vanishes (Delaunay, 1933). The seven sums of subsets of an obtuse superbase are the Voronoi vectors of the lattice, and the pattern of vanishing dot products fixes its Voronoi type, one of five parallelohedra (Voronoi, 1908; Conway & Sloane, 1992). Kurlin (2022) listed all obtuse superbases of each Voronoi type and built from their conorms the root invariant, a complete and continuous invariant of the lattice up to isometry. Delone's 24 sorts classify lattices by the zeros and equalities among the six dot products (Delone *et al.*, 1975; Engel *et al.*, 2004; Andrews & Bernstein, 2023). Andrews and Bernstein treat the six numbers as a point of S^6, where the reduced presentations of a lattice fill a cone whose boundaries are the vanishing components and where boundary transformations carry one reduced presentation to another (Andrews *et al.*, 2019*b*, 2023, Section 6.3; Andrews & Bernstein, 2014); this geometry underlies the distances CS6Dist and NCDist and the Bravais-type determination of SELLA. Oishi-

Tomiyasu (2012, 2024) gives error-stable Bravais-lattice determination in Selling coordinates with a provably sufficient list of basis transforms. The instability of the Niggli cell (Andrews *et al.*, 1980; Andrews & Bernstein, 1988, 2014; Gruber, 1989) is the same multiplicity seen from the Niggli side, and the operators between two settings of one lattice form a coset of its symmetries (Flack, 1987).

Given this, what the paper adds is four things, each a statement with a guarantee.

1. *The matrices between reduced cells, explicitly.* The boundary transformations of a lattice compose to finitely many integer matrices between its reduced cells. We give them as one universal set $\mathcal{U}$ of $6096 = 48 \times 127$ matrices with entries up to ± 3: every matrix between two reduced cells of any lattice lies in $\mathcal{U}$, every element of $\mathcal{U}$ is needed for some lattice, and which elements occur depends only on which conorms vanish. The proof is a short lemma and an exhaustive enumeration, checked in Lean (Section 3.3, Lemma A2). The entries are bounded exactly, Voronoi type by Voronoi type and labelling by labelling (Table 4).

2. *An enumeration under a tolerance.* For a measured cell, the superbases that are reduced for some metric within a strain η are finitely many and are enumerated (Lemma A6), and reindexing a new cell against a reference widened this way returns exactly the proper operators within a price band (Lemma A7). Whether $\mathcal{U}$ alone suffices for a given cell is certified by a radius computed for that cell (Remark A8); for all but a thin tail of cells it does.

3. *Prices that do not depend on the basis.* A candidate is priced by the affine-invariant distance between the two metric tensors after the best common scale. The candidates lie between two cosets of the reference's near-symmetries, the near-symmetries are bounded by the symmetries of the nearest symmetric lattice (Proposition A10), and a pseudo-twofold costs exactly twice its distance to the metrics it would leave invariant (Proposition A9).

4. *A measurement of what the cell cannot decide.* The price between deposited cells of one crystal form sets the scale of the band, and comparing it with the prices of the pseudo-symmetries in the PDB bounds how often the cell alone leaves an alternative indexing open (Section 5.3).

Reindexing as described here returns exactly the proper operators whose price is within the band. Two consequences follow. The true operator is among them whenever its own price is within the band. And no method that uses the two cells alone can return fewer while keeping that guarantee, since any operator within the band would be the true one if the two cells differed by its price; the distinct indexings are then the classes of these operators under the Laue group (Section 5.2). The guarantee holds within a stated scope: primitive relations of determinant ± 1, strains below 25%, enumerations of at most 5000 lattice vectors, and membership decided with one consistent relative tolerance of 10^{-9}, as Grosse-Kunstleve *et al.* (2004) did for Niggli reduction (Section 6). The searches in use find the true operator in ordinary cases, and nothing here is offered as a faster or more successful search; what is new is the guarantee and the price of each alternative. Everything is implemented in the open-source package agentsg (Zwart, 2026), which also reproduces every number in the paper.[1]

Section 2 collects the classical definitions, Section 3 the theory with a table of what is proved, Section 4 the methods and tests, and Section 5 the results, ending with the PDB application; the results the theory rests on are summarised in Appendix A, and their proofs, catalogues, further tables and the computational checks are in the supporting information.

2. Background

2.1. Superbases, conorms and the metric tensor

Selling reduction works with four vectors: the cell vectors a, b, c and $d = -a - b - c$. The four form a *superbase*, and any three of them are a basis of the lattice. The six negated pair products

$$p_{da} = -d \cdot a, \quad p_{db} = -d \cdot b, \quad p_{dc} = -d \cdot c,$$

$$p_{ab} = -a \cdot b, \quad p_{ac} = -a \cdot c, \quad p_{bc} = -b \cdot c$$

are the *conorms* (Selling, 1874; Conway & Sloane, 1992). Conorm vectors are always written in this slot order, da, db, dc, ab, ac, bc; the Selling parameters of S^6 are the same numbers with the opposite sign, in another order. A superbase is *obtuse*, which is to say Selling-reduced, when all six conorms are non-negative. The metric tensor in the basis (a, b, c) is linear in the conorms,

$$\mathbf{G} = \begin{pmatrix} p_{ab} + p_{ac} + p_{da} & -p_{ab} & -p_{ac} \\ -p_{ab} & p_{ab} + p_{bc} + p_{db} & -p_{bc} \\ -p_{ac} & -p_{bc} & p_{ac} + p_{bc} + p_{dc} \end{pmatrix}, \tag{1}$$

so the conorms are six cell parameters. A zero conorm is a right angle between two of the four vectors, and the set of slots that vanish is the *zero pattern*. The conorm vectors with $p \geq 0$ and a given zero pattern form a boundary of the cone of reduced conorms, in S^6 a boundary of the region of reduced presentations (Andrews & Bernstein, 2014); it describes lattices only if its non-zero slots connect the four vectors, and 38 zero patterns do. Up to relabelling they give the five Voronoi types (Voronoi, 1908; Conway & Sloane, 1992; Kurlin, 2022): no zero (Voronoi cell a truncated octahedron), one zero (a hexarhombic dodecahedron), two zeros on disjoint pairs (a rhombic dodecahedron), two zeros sharing a vector (a hexagonal prism), and three zeros

[1] This work used AI models (Claude Fable 5.1 and Claude Sonnet 5.5, Anthropic) extensively for coding, proofs, insight and writing. To ensure the correctness of the results, every statement and proof that can be computed has a computational check: a machine-checked proof in Lean, or a program that verifies it by exhaustive enumeration or by spot checks on seeded samples. The checks are listed, each under the number of the statement it verifies, in the supporting information (Section S7).

(a cuboid), the last in two placements, a triangle on a, b, c and a path through d, which describe the same lattices in different cells. The seven sums of subsets of a superbase are its *partial sums*; for an obtuse superbase they are the Voronoi vectors of the lattice, one in each class of lattice vectors modulo twice the lattice other than the doubled vectors themselves (Conway & Sloane, 1992). We call these classes *parity classes*: written as $x = n_1a + n_2b + n_3c$, a vector's class is given by which of the coefficients n_i are even and which odd (Appendix A).

2.2. Selling steps and relabellings

Two kinds of integer operator act on a superbase.

Selling steps. The Selling step at the conorm p_{ab} replaces a by $-a$ and adds the old a to the two other vectors, $(a, b, c, d) \mapsto (-a,\ b,\ c+a,\ d+a)$. The new conorms are $p'_{ab} = -p_{ab}$, $p'_{ac} = p_{da} + p_{ab}$, $p'_{da} = p_{ac} + p_{ab}$, $p'_{bc} = p_{bc} + p_{ab}$, $p'_{db} = p_{db} + p_{ab}$ and $p'_{dc} = p_{dc} - p_{ab}$. Selling reduction applies steps at negative conorms until none remain. At a conorm that is exactly zero the step exchanges the slots *ac* and *da* and carries one obtuse superbase to another of the same lattice: it is a *boundary transformation* in the sense of Andrews *et al.* (2023, Section 6.3), and it is why the reduced cell of a lattice with a right angle is not unique. In the basis (a, b, c) each step is a fixed 3×3 integer matrix of determinant -1 that is its own inverse (Table 1). A word of steps such as *ab.bc.db* is applied left to right, each step being read in the labelling current at that point, and its matrix is $\mathbf{F}_{ab}\mathbf{F}_{bc}\mathbf{F}_{db}$, columns being the new basis vectors in the old basis.

Table 1
The six Selling steps as 3×3 integer matrices in the basis (a, b, c), $d = -a - b - c$; columns are the new basis vectors in the old basis. Each has determinant -1 and squares to the identity. For the steps at *da*, *db* and *dc* the vector d is negated and the new a, b, c are listed.

step at	new basis	matrix
ab	$(-a,\ b,\ a+c)$	$\begin{pmatrix} -1 & 0 & 1 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$
ac	$(-a,\ a+b,\ c)$	$\begin{pmatrix} -1 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$
bc	$(a+b,\ -b,\ c)$	$\begin{pmatrix} 1 & 0 & 0 \\ 1 & -1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$
da	$(a,\ -a-c,\ -a-b)$	$\begin{pmatrix} 1 & -1 & -1 \\ 0 & 0 & -1 \\ 0 & -1 & 0 \end{pmatrix}$
db	$(-b-c,\ b,\ -a-b)$	$\begin{pmatrix} 0 & 0 & -1 \\ -1 & 1 & -1 \\ -1 & 0 & 0 \end{pmatrix}$
dc	$(-b-c,\ -a-c,\ c)$	$\begin{pmatrix} 0 & -1 & 0 \\ -1 & 0 & 0 \\ -1 & -1 & 1 \end{pmatrix}$

Relabellings. A superbase is an unordered set of four vectors up to a common sign, so one reduced cell can be written in $4! \times 2 = 48$ ways: a *labelling* says which three of the four vectors, in which order and with which common sign, are called a, b, c. The 48 relabellings change which vectors are called a, b, c and nothing else; on the six conorms they act as the 24 permutations that Andrews *et al.* (2019*b*) call the reflections of S^6, the common sign acting trivially. Negating a single basis vector is not a relabelling: it does not preserve $a+b+c+d = 0$.

2.3. Delone sorts and the root invariant

Delone (published as Delaunay, 1933) classified lattices into 24 sorts by the pattern of zeros and equalities among the reduced conorms. Each sort lies in one Bravais type, and some Bravais types comprise several sorts (Delone *et al.*, 1975; Burzlaff & Zimmermann, 1985, 2006). We use the symbols C1 to A3 and the patterns of Andrews & Bernstein (2023) and SELLA (Andrews *et al.*, 2023), read in the slot order of Section 2.1; the correspondence with their slot order is given in the supporting information. The sorts are the Selling counterpart of the 44 Niggli lattice characters (Niggli, 1928; de Wolff & Gruber, 1991). They split a Bravais type by the shape of the reduced cell, not by its symmetry. A zero conorm restricts the metric without adding a symmetry: one right angle between two superbase vectors gives no twofold, so the triclinic sorts A2 and A3, with one and two zeros, are as triclinic as A1.

All obtuse superbases of a lattice carry the same six conorm values as a multiset, since a boundary transformation only exchanges two slots. Which slots hold which value is not invariant: in the monoclinic example the flip at *ab* moves p_{ac} into the slot *da*. Kurlin (2022) normalises the square roots of the conorms, by a rule that depends on the Voronoi type, to the root invariant, a complete isometry invariant of the lattice. The invariant identifies a lattice; the closure records how its reduced cells are related.

2.4. Terms

Table 2 collects the terms used in the rest of the paper, with the section in which each is defined. Strains and prices are given in per cent throughout; a band of 7.4% is 0.074 in the formulas. A strain measures a change of metric and a price a distance between two metrics (Section 3.4). We write η for a strain and ϵ for a price band; a price band ϵ needs the reference widened at the strain $\eta = e^{\epsilon} - 1$ (Lemma A7), 7.7% for the default band of 7.4%.

3. Theory

3.1. The Selling closure

A lattice has no single preferred reduced cell. It has a finite set of obtuse superbases, and two kinds of integer matrix relate them: symmetries, which map a setting to itself, and reindexings, which map one setting to another. The *Selling closure* $\mathcal{C}(\mathbf{G})$ of a lattice collects both: the metric tensors of its obtuse superbases in all their labellings, and, from $\mathbf{G}$ to $\mathbf{G}'$, the unimodular matrices $\mathbf{P}$ with $\mathbf{P}^{\mathsf{T}}\mathbf{G}\mathbf{P} = \mathbf{G}'$. Labellings related by a symmetry give the same metric tensor and count once. In categorical terms the closure is a connected groupoid (Brandt, 1927; Higgins, 1971), the part on one lattice's reduced cells of the action groupoid of $GL_3(\mathbb{Z})$ on positive definite matrices (Weinstein, 1996). A connected groupoid is equivalent to the group of its loops at any one object, here the lattice point group, so as a structure it carries nothing beyond the point group and its

Table 2
Terms and symbols.

term	meaning	defined
superbase, conorm, obtuse	four lattice vectors summing to zero; the six negated pair products p_{ij}; all six conorms ≥ 0 (Selling-reduced)	2.1
reduced cell	an obtuse superbase, in some labelling	2.1
zero pattern, Voronoi type	the slots whose conorms vanish; the five types of Voronoi cell they give up to relabelling	2.1
partial sum, parity class	the sum of some of the four vectors (a Voronoi vector when the superbase is obtuse); a class of lattice vectors modulo twice the lattice	2.1
Selling step, boundary transformation	the operators of Table 1; a Selling step at a zero conorm	2.2
labelling, relabelling	which three vectors of a superbase, in which order and common sign, are a, b, c; the 48 changes of labelling	2.2
sort	one of Delone's 24 classes of lattices by zeros and equalities among the conorms	2.3
closure $\mathcal{C}(\mathbf{G})$, shape	the obtuse superbases of a lattice in all their labellings, with the matrices between them; superbases whose conorms agree up to a relabelling	3.1
$\mathcal{U}$	the 6096 integer matrices between two reduced cells of some lattice	3.3
strain η; geodesic distance D	$\|\mathbf{G}^{-1/2}\boldsymbol{\Delta}\mathbf{G}\,\mathbf{G}^{-1/2}\|_F$, equation (2); the affine-invariant distance, equation (3)	3.4
price, band ϵ	the geodesic distance, after the best common scale, between a metric and the one an operator predicts; the price within which an operator is accepted	3.4, 3.4
depth, widened set $\mathcal{S}_\eta$	the strain to the nearest metric that makes a superbase obtuse; the superbases of depth at most $\eta = e^\epsilon - 1$	3.4
first layer, radius r_1	the 127 superbases $\mathcal{U}F_0$ that can be obtuse together with F_0; the strain at which a superbase outside it first becomes obtuse	3.4
reference, new cell, candidate	the stored cell, the cell to be reindexed against it, and an operator between them within the band	4.2
near-symmetry H_ϵ, pseudo-symmetry	the proper unimodular self-maps of price at most ϵ; a symmetry of a nearby lattice of higher symmetry	3.4
indexing class	a class of the Laue group (point group for anomalous data) among the candidates	3.4
obliquity δ, non-isomorphism	Le Page's angle between a lattice row and the nearest reciprocal row; the price between two crystals of one form	3.4, 4.3

set of objects; the content of this paper is computational: the finite set of reduced cells of each lattice, the one set $\mathcal{U}$ that holds every matrix between them, and the dependence of those matrices on the zero pattern alone. The matrices from a metric tensor to itself are the lattice point group in that basis (Himes & Mighell, 1987); those between two are one coset of it (Flack, 1987). They are known in advance: a lattice has $n = 1, 2, 3, 6$ or 16 obtuse superbases (unlabelled) according to its Voronoi type and zero pattern (Kurlin, 2022; Lemma A2), every matrix of the closure is a word of boundary transformations followed by a relabelling (Lemma A3), and every matrix lies in one fixed set $\mathcal{U}$ (Section 3.3). The symmetries that merely relabel a reduced cell leave it in place; the superbases they reach from it are those of the same *shape*, and the order of the point group is the number of relabelling symmetries times the number of superbases of that shape. Every other symmetry carries the reduced cell to a different reduced cell. A reduced cell here is always an obtuse superbase; a Niggli cell of type II (no acute angle) is one, and a Niggli cell of type I (three acute angles) is not.

3.2. One monoclinic cell

Take $a = 30$, $b = 40$, $c = 50\,\text{Å}$, $\beta = 100°$. Its reduced superbase is $M_0 = (d, a, b, c)$ with conorms

$$p_{da} = 639.5, \quad p_{db} = 1600.0, \quad p_{dc} = 2239.5,$$

$$p_{ab} = 0, \quad p_{ac} = 260.5, \quad p_{bc} = 0\ \text{Å}^2.$$

The two zeros are the right angles α and γ at the twofold axis b, present in every primitive monoclinic lattice. The closure has six superbases (Table 3). Each has two zero conorms, each boundary transformation leads to a neighbour, and the six form a ring M_0–M_1–M_5–M_4–M_3–M_2–M_0 (Fig. 1). They fall into three shapes. $\{M_0, M_4\}$ is 30, 40, 50 Å, $\beta = 100°$; $\{M_1, M_3\}$ is 30, 40, 53.7 Å, $\beta = 113.4°$, the cell on a and $a+c$; $\{M_2, M_5\}$ is 50, 40, 53.7 Å, $\beta = 146.6°$, the third cell choice, on c and $a + c$. These are the three cell choices of *International Tables* (Hahn, 2006, Section 2.2.16).

Table 3
The Selling closure of the monoclinic cell $a = 30$, $b = 40$, $c = 50\,\text{Å}$, $\beta = 100°$: six obtuse superbases, each up to the order and common sign of its vectors, the boundary transformations that connect them, and where the four symmetries of the point group $2/m$ carry M_0.

superbase	four vectors	boundary transformations lead to	symmetries carry M_0
M_0	$-a-b-c,\ a,\ b,\ c$	M_1, M_2	1, $\bar{1}$ (relabellings c
M_1	$-b-c,\ -a,\ b,\ a+c$	M_0, M_5	
M_2	$-a-c,\ a+b,\ -b,\ c$	M_0, M_3	
M_3	$a+c,\ b-c,\ -b,\ -a$	M_2, M_4	
M_4	$a-b+c,\ -c,\ b,\ -a$	M_3, M_5	2 along b, m norm
M_5	$a-b,\ c,\ b,\ -a-c$	M_1, M_4	

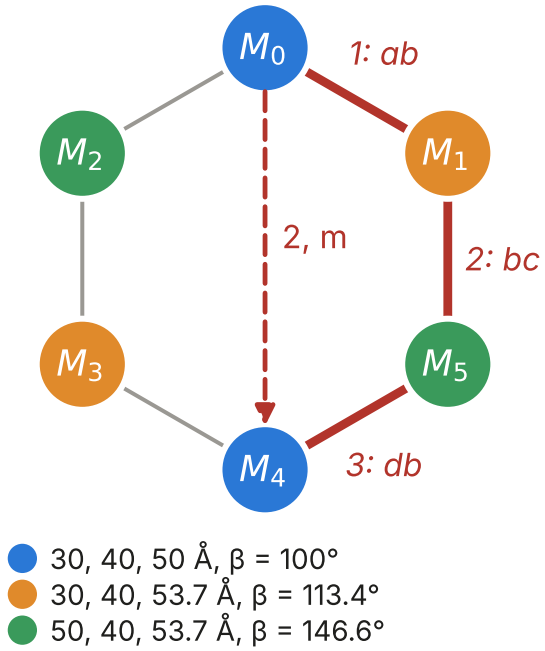


Figure 1
The closure of the 100° cell: six superbases on a ring, joined by boundary transformations at the two zero conorms and coloured by cell shape. The red path is the word *ab.bc.db* from M_0 to M_4, each step read in the labelling current at that point. The twofold and the mirror carry the vectors of M_0 to those of M_4 along this path, followed by a relabelling; they relabel no superbase on the way, so they do not show in the reduced cell.

The point group is 2/*m*, and the last column of Table 3 shows how its four operations act on the six superbases. The identity and the inversion relabel M_0. The twofold sends $(a, b, c) \mapsto (-a, b, -c)$ and so carries the four vectors $(-a-b-c, a, b, c)$ to $(a-b+c, -a, b, -c)$, which are the vectors of M_4, three boundary transformations away. As a matrix of the closure each symmetry maps the metric tensor of M_0 to itself; as a word, the twofold is the steps *ab*, *bc*, *db* — the path $M_0 \to M_1 \to M_5 \to M_4$ — followed by the relabelling that renames the vectors of M_4 so that its metric tensor is that of M_0, and the mirror is the same steps with that relabelling composed with the change of common sign. Half of the point group therefore does not show in the reduced cell: two symmetries relabel M_0, and M_0 and M_4 have the same shape, as the count $4 = 2 \times 2$ requires. A search over lattice directions, such as Le Page's, finds the twofold either way. Two measurements that reduce to M_0 and to M_1 report different cell parameters and are related by no symmetry, only by the boundary transformation at *ab*.

3.3. The universal operator set and the size of its entries

The matrices can be characterised with no lattice at all. An integer matrix **W** acts on the six conorms by a fixed integer 6×6 matrix, and **W** maps a reduced cell to a reduced cell for every lattice with a given zero pattern exactly when that matrix maps the corresponding boundary of the cone into the cone of non-negative conorms, a sign condition on its entries (supporting information, Section S3). The union over the 38 zero patterns is the set $\mathcal{U}$. It has $6096 = 48 \times 127$ elements with entries in $\{-3, \ldots, 3\}$ and is closed under inversion. By Lemma A2 and the exhaustive enumeration of its window (Appendix A, the enumeration after Lemma A2), $\mathcal{U}$ contains every matrix between two reduced cells of any lattice; this, and the converse that every element of $\mathcal{U}$ is such a matrix for every lattice with some zero pattern, are checked in Lean (Section 3.5). The cone of non-negative conorms is Voronoi's domain of the principal perfect form (Voronoi, 1908; Schürmann, 2009), and the 127 superbases $\mathcal{U}F_0$ of a reduced superbase F_0 are the domains that meet it within the positive definite metrics: all the superbases that can be obtuse together with F_0 for some metric, the *first layer*. The count follows by hand: the reduced cell itself, one superbase for each of the 6 single zeros, 3 new ones for each of the 12 pairs of zeros sharing a vector and for each of the 4 triangles of three zeros, 6 for each of the 12 paths, and none from two zeros on disjoint pairs, $1 + 6 + 36 + 12 + 72 = 127$ (Appendix A); we have not found it stated before. The first layer carries the completeness statements of Section 3.4.

Which matrices occur depends only on the zero pattern. A superbase is obtuse exactly when each of its partial sums is shortest in its parity class, and by Lemma A1, written relative to F_0, that holds exactly when certain integers vanish on the slots where F_0's conorms are non-zero. A superbase obtuse for one metric with a given zero pattern is therefore obtuse for every metric with that pattern (checked in Lean), one generic metric per pattern describes all its lattices, and the 38 patterns cover all lattices. The size of the entries depends on a convention, since relabelling either cell changes them: every element of $\mathcal{U}$ is a relabelling, times a matrix with entries in $\{-1, 0, 1\}$, times a relabelling. With the convention of Table 4, by enumeration (Lemma A2): with no zero conorm, one zero, or two zeros on disjoint pairs, every operator has entries in $\{-1, 0, 1\}$; with two zeros sharing a vector or with three zeros, the entries reach 2 when a zero sits on a slot of $d = -a-b-c$, and 3 for a path of three zeros with two on d, such as conorms $(0, 0, 3, 2, 0, 5)$ with the obtuse basis $(-3a-2b-c,\ a,\ a+b)$. Lemma A2 alone allows larger entries; writing $d = -a-b-c$ is what turns a superbase coefficient of 2 into an entry of 3. A relabelling of the reduced cell that moves its zeros off d brings every operator into $\{-1, 0, 1\}$ wherever such a labelling exists, for one zero, two zeros sharing a vector and a triangle of three. A path of three zeros touches all four vectors, so no labelling keeps it off d, and in the best labelling of the first cell an operator still needs an entry of 2; relabelling both cells removes it.

Table 4
The largest entry of an operator from a reduced cell to another reduced cell of the same lattice, by zero pattern and by whether a zero conorm sits on a slot of d, over all 38 patterns: the first cell is taken as reduction returns it, in its own basis, and the second in its worst labelling. By enumeration (Lemma A2).

zero conorms	zeros avoid d	a zero on d
none	1	
one	1	1
two on disjoint pairs (O3, cF, R3, M3, A3)	—	1
two sharing a vector (mP, oS, hP)	1	2
three, a triangle (oP, tP, cP from an orthogonal cell)	1	2
three, a path (the same lattices from a face-diagonal cell)	—	2 or 3

3.4. Measured cells

A measured cell is rarely exactly on a zero conorm. It is when symmetry has been imposed: a cell refined in its space group, or, in serial crystallography, a fixed reference cell assigned to every crystal (White *et al.*, 2012; Hattne *et al.*, 2014). A cell refined in a lower symmetry, or indexed without constraints, measures a right angle as, say, 90.25°, and the conorm is a small number of either sign. The flips that joined the superbases of Table 3 then lead to superbases with a slightly negative conorm, Selling reduction rejects them, and the exact closure of the measured metric has one superbase instead of six. To compare such a cell with one measured or constrained differently, the closure has to be widened.

Two measures of mismatch. Two measures are used, both the same in every basis. The *strain* of a perturbation $\Delta\mathbf{G}$ of a metric $\mathbf{G}$ is

$$\|\Delta\mathbf{G}\|_{\mathbf{G}} = \|\mathbf{G}^{-1/2}\,\Delta\mathbf{G}\,\mathbf{G}^{-1/2}\|_F. \tag{2}$$

It is unchanged by a change of basis, since $\mathbf{G} \mapsto \mathbf{P}^{\mathsf{T}}\mathbf{G}\mathbf{P}$ and $\Delta\mathbf{G} \mapsto \mathbf{P}^{\mathsf{T}}\Delta\mathbf{G}\mathbf{P}$ leave it as it is, and by a common scale of the cell. A perturbation of strain η changes the dot product of any two lattice vectors by at most $\eta|u||v|$ (Proposition A4). The strain measures how far a metric must move to make a superbase obtuse, a convex problem in its six conorms. The *geodesic distance* between two metrics is

$$D(\mathbf{G}_1, \mathbf{G}_2) = \Big(\sum_i \ln^2 \boldsymbol{\lambda}_i\Big)^{1/2}, \tag{3}$$

with λ_i the eigenvalues of $\mathbf{G}_1^{-1/2}\mathbf{G}_2\mathbf{G}_1^{-1/2}$, the affine-invariant distance on positive definite matrices (Moakher, 2005; Pennec *et al.*, 2006; Bhatia, 2007). It agrees with the strain to first order, it is a true distance (symmetric, with the triangle inequality), and every change of basis is an isometry for it. It measures how far apart two cells are, and it prices operators. Both are finite-strain measures of the deformation that carries one lattice onto the other: the strain is twice the norm of the Green–Lagrange strain tensor computed from the two sets of cell parameters (Schlenker *et al.*, 1978), as in studies of thermal expansion and compressibility (Cliffe & Goodwin, 2012), and the geodesic distance is twice the norm of the logarithmic (Hencky) strain, which is itself a geodesic distance (Neff *et al.*, 2016).

Table 5 converts the classical tolerances. For Le Page's twofold criterion (Le Page, 1982) the conversion is exact: a twofold whose obliquity, the angle between a lattice row and the nearest reciprocal row, is δ sits at a geodesic distance $\sqrt{2}\,\operatorname{artanh}(\sin\delta)$ from the metrics it would leave invariant, and its price is twice that (Proposition A9). One degree of obliquity is a price of 4.9%, and Le Page's 3° a price of 14.8% (Table 5), for every twofold separately.

The widened set. The *widened closure* of a metric $\mathbf{G}$ at strain η is

$$\mathcal{S}_\eta(\mathbf{G}) = \big\{\, F \;:\; F \text{ is obtuse for some metric within strain } \eta \text{ of } \mathbf{G} \,\big\}. \tag{4}$$

Its superbases are the reduced cells of every lattice consistent with the measurement at that tolerance. Each superbase has a *depth*, the strain from $\mathbf{G}$ to the nearest metric that makes it obtuse, and belongs to $\mathcal{S}_\eta(\mathbf{G})$ when its depth is at most η. Equation (4) names no starting cell, order of operations or basis, so $\mathcal{S}_\eta(\mathbf{G})$ is the same whoever computes it. Each operator $\mathbf{P}$ from a metric $\mathbf{G}_1$ to a metric $\mathbf{G}_2$ has a *price*, the geodesic distance between $\mathbf{G}_2$ and $s\mathbf{P}^{\mathsf{T}}\mathbf{G}_1\mathbf{P}$ minimised over a common scale s, which removes the trace of the logarithm in equation (3). The price is zero for every matrix of an exact closure, and for a unimodular self-map the scale changes nothing. Since $\mathbf{P}$ and $-\mathbf{P}$ predict the same metric, only proper operators are returned, which loses nothing.

It is finite. The seven partial sums of an obtuse superbase are shortest vectors of their parity classes (Section 2.1). With μ_c the squared length of the shortest vector of class c, every vector of a superbase in $\mathcal{S}_\eta(\mathbf{G})$ lies in a ball of squared radius $\frac{1+\eta}{1-\eta}\mu_c$ (Lemma A6), so the candidates for $\mathcal{S}_\eta(\mathbf{G})$ are finitely many and are listed with no window of coordinates to assume. Finiteness is not smallness: for needles and plates at large strains the balls hold many vectors, and the implementation refuses more than 5000 (Section 5.1). Membership is then decided by computing each candidate's depth, a convex problem, in floating point, with one relative tolerance of 10^{-9} applied consistently, as Grosse-Kunstleve *et al.* (2004) did for the comparisons of Niggli reduction; only a superbase whose depth lies within that tolerance of η could be decided either way.

Widening one side suffices. If an operator $\mathbf{P}$ from a reference to a new cell has price at most ϵ, the new cell's reduced superbase, carried back by $\mathbf{P}$, lies in the reference's widened closure at strain $\eta = e^{\epsilon} - 1$ (Lemma A7). So only the reference is widened, and the new cell is reduced once.

When the 6096 matrices suffice. When the widened set lies in the first layer $\mathcal{U}F_0$, the matrices of $\mathcal{U}$ between the two Selling-reduced cells already contain every operator within the band (Remark A8). The enumeration reports whether it does, and a radius computed for each cell certifies it in advance: $r_1(\mathbf{G})$, the strain to the nearest metric for which a superbase outside the first layer is obtuse, searched up to 25%, certifies every band ϵ with $e^{\epsilon} - 1 < r_1$. Beyond r_1 the fixed set is not enough. Against the $100 \times 100.1 \times 5$ Å plate of Table 8, whose r_1 is 5%, a new cell whose true operator costs 6.1% is found by the enumeration and missed by the 6096 matrices (Section 5.1): $\mathcal{U}$ is complete for exact cells, and for measured cells only below r_1.

Table 5
Classical tolerances as strains and prices. The first two rows give, to first order, the smallest strain of such a change, reached for an orthogonal pair (first row) or an axis normal to the other two (second row); the last two rows are exact for a twofold along a lattice row, in every cell (Proposition A9).

tolerance	strain or price	value
angle between two vectors changed by δ, edges fixed	at least $\sqrt{2}\sin\delta$	7.4% at 3°
edge changed by a fraction f, angles fixed (an axis normal to the other two)	at least $2f$	10% at 5%
twofold of obliquity δ: geodesic distance to the metrics it fixes	$\sqrt{2}\,\mathrm{artanh}(\sin\delta)$	7.4% at 3°
twofold of obliquity δ: its price	$2\sqrt{2}\,\mathrm{artanh}(\sin\delta)$, 4.9% per degree	14.8% at 3°

Near-symmetries and indexing classes. The self-maps of a measured metric of price at most ϵ form a finite set that contains the point group, and two of them compose to one of price at most the sum; they need not compose within the band, so the widened set has no group of symmetries of its own. Let H_ϵ be the proper ones and $p_0 \le \epsilon$ the price of the true operator $\mathbf{P}_0$. Every candidate is $h\mathbf{P}_0$ with $h \in H_{\epsilon+p_0}$, and every $h \in H_{\epsilon-p_0}$ gives one; the near-symmetries are bounded on both sides by the symmetries of the nearest symmetric lattice (Proposition A10). Two candidates that differ on the left by an element of the reference's Laue group $\mathcal{L}$ index the crystal the same way, so the distinct indexings are the classes of the proper part $\mathcal{L}^+$ among the candidates; for anomalous data from an acentric crystal the point group replaces the Laue group. The price orders the classes but does not pick the true member, which is a matter for the intensities.

The price of pseudo-symmetry. A symmetry g of a nearby lattice enters H_ϵ when its price is at most ϵ. For a twofold the price is exactly twice the geodesic distance from the measured metric to the metrics V_g that g leaves invariant, which gives the exact conversion from Le Page's obliquity of Table 5. For rotations of higher order there is no such exact conversion: the price is at most twice the distance, and to first order near V_g the ratio is $\sqrt{3}$ for a threefold, between $\sqrt{2}$ and 2 for a fourfold and between 1 and $\sqrt{3}$ for a sixfold (Proposition A9), so the rule of about 5% per degree applies to twofolds only.

3.5. What is proved

Table 6 separates what is proved from what rests on an exhaustive computation and what is checked. "Machine-checked" marks statements proved in the proof assistant Lean 4 with Mathlib (de Moura & Ullrich, 2021; mathlib Community, 2020). Where a theorem rests on a finite computation, as the completeness of $\mathcal{U}$ does, the computation is evaluated by compiled code inside Lean, which adds the Lean compiler to the kernel as something trusted; the proof then consists of a bound proved in the kernel and finite checks evaluated this way. The formal statements, the versions and the axioms each theorem rests on are in the supporting information (Section S7, Table S12), and the enumeration that completes Lemma A2 was also run independently. The affine-invariant distance enters the formal statements only through the properties stated in Section 3.4 and, for Proposition A9(i), the existence and uniqueness of midpoints. The proofs rest on Selling's identity, the characterisation of obtuse superbases by Voronoi vectors (Conway & Sloane, 1992; Appendix A) and Selling reduction. Exact closures and the closures of the 24 sorts were checked against brute force over bases with entries in $[-2,2]$, a window that contains a labelling of every member of every exact closure (by the enumeration that completes Lemma A2), and widened sets against the ball enumeration of Lemma A6.

4. Methods

All methods are implemented in agentsg (Zwart, 2026), version 0.5.0; the test cases and the PDB analysis are in the same repository (supporting information, Sections S2 and S7).

4.1. Computing the widened set

For a measured cell and a strain $\eta < \frac{1}{4}$: Selling-reduce the cell to a superbase F_0; enumerate the lattice vectors in the seven balls of Lemma A6 and form the superbases they make; keep those of depth at most η, with each one's depth and the operator $\mathbf{W}$ from F_0 to it. The list of candidates is exact for every cell, and membership is decided with the relative tolerance of 10^{-9} of Section 3.4. Walking from the reduced cell, taking a Selling step wherever a conorm is within a band of zero, is not exhaustive: a member can be reachable only through a non-member, when two small conorms lie at, say, 1.2 and 0.5 times the band and only the second of two steps returns into it. The benchmark contains such members, and walks that miss them are compared in the supporting information (Section S2, Table S3).

Ties take no band. Every sort lies on one zero pattern, and every symmetry of every lattice on that pattern is one of its matrices (Lemma A3), so a near-equality — the $a \approx b$ of a pseudo-tetragonal cell — never adds a superbase: the operator that would be its symmetry is already enumerated, and its price decides whether it is accepted.

4.2. Reindexing

A reference R is stored as its widened closure, the operators $\mathbf{W}_k$ from its reduced cell to the members. A new cell N is reduced once, to a basis B_N in its own coordinates; B_R is the reduction matrix of the reference. For each stored $\mathbf{W}_k$ and each of the 48 relabellings $\mathbf{Q}$ the candidate operator from the reference's input basis to the new cell's is

$$\mathbf{P} = B_R \mathbf{W}_k \mathbf{Q}\, B_N^{-1}, \tag{5}$$

so that $(a_N\, b_N\, c_N) = (a_R\, b_R\, c_R)\mathbf{P}$, $\mathbf{G}_N \approx \mathbf{P}^{\mathsf{T}}\mathbf{G}_R\mathbf{P}$, and Miller indices transform as $h_N = \mathbf{P}^{\mathsf{T}} h_R$: for the Selling step at ab of Table 1, $\mathbf{P} = \mathbf{F}_{ab}$, the new cell is $(-a,\ b,\ a+c)$, and the reflection $(1,0,0)$ of the reference is $(-1,0,1)$ in the new cell. Candidates of price at most the band ϵ are returned; the reference

Table 6
What is proved, what rests on an exhaustive computation, and what is checked or refuted.

statement	status
every matrix of a closure is a word of boundary transformations and a relabelling	proved (Lemma A3)
the members of each closure lie in a finite window	proved (Lemma A2); machine-checked, in the basis (a, b, c) with coordinates in $[-7, 7]$
$\mathcal{U}$ contains every matrix between reduced cells of any lattice; every element of $\mathcal{U}$ is such a matrix for some zero pattern; the matrices depend only on the zero pattern; the entry sizes by zero pattern	computer-assisted: the window enumerated over all 38 zero patterns; the first three machine-checked against the deposited table, a bound proved in the kernel with finite checks evaluated by compiled code
the 24 sorts	a definition and an enumeration, matched to Delone's patterns by computation (Definition A5)
the strain bounds every dot product	proved (Proposition A4)
the widened set is finite	proved (Lemma A6); its candidates listed exactly, membership decided by a floating-point depth with one relative tolerance of 10^{-9}
in the first layer, the 6096 matrices are complete	proved (Remark A8, from Lemma A7, the completeness of $\mathcal{U}$ and its closure under relabelling), the last two machine-checked; the radii of Table 8 certified in exact arithmetic for the stored conorms, those of the archive statistics computed in floating point, up to 25%
widening the reference alone suffices, at $\eta = e^{\epsilon} - 1$	proved (Lemma A7)
a twofold's price is twice its distance to the metrics it fixes, and $2\sqrt{2}\,\operatorname{artanh}(\sin\delta)$	proved, exactly (Proposition A9); the bound of part (i) machine-checked for any isometric action, and the fixed midpoint of a twofold given that midpoints exist and are unique
the ratios of price to distance for orders 3, 4, 6	proved in the strain at a symmetric metric; first order for the price (Proposition A9)
a type II Niggli cell is obtuse, with its small conorms off d below a stated size, near two or three right angles; a type I metric near a three-zero lattice reduces to a path through d from any basis	proved (Proposition A11), parts (i) and (ii) machine-checked; type I near two right angles, and reduction from other bases, observed
two Niggli cells of metrics with one zero pattern differ by a matrix in $\{-1, 0, 1\}$, and so do those of two measurements within ρ of one lattice; the Niggli vectors are signed partial sums of every obtuse superbase, so one fixed list of 1 344 matrices (672 up to sign) leads from every reduced cell to the Niggli cell	computer-assisted (Proposition A12, Corollary A13): every basis and superbase over all 38 zero patterns, exclusions by exact rational certificates; ρ computed numerically, to first order in the strain
candidates lie between two cosets of self-maps; H_ϵ is a group in a gap of the prices; the near-symmetries lie between the point group of the nearest symmetric lattice and its self-maps within $\epsilon + 2\nu$	proved, exactly (Proposition A10); machine-checked for any isometric action
the first layer of 127 superbases (also counted by hand; supporting information, Section S12) and the second of 2755; the largest entries after relabelling; Lemma S4.4 and so Theorem S4.5 of an absolute band on the conorms (supporting information, Section S4)	computer-assisted enumeration, the layers also recomputed independently

is widened at $e^{\epsilon} - 1$, so the true operator is among them whenever its price is at most ϵ (Lemma A7). The candidates need not be built one by one. The two input bases drop out of every price, so the reference's members are tabulated once as the metrics they predict, 24 per member; every candidate has the same best scale; a necessary condition on the dot products of the new cell's reduced superbase, from Lemma A7 and Proposition A4, discards most of them before any price is computed; and the exact operator is built only for those within the band. This returns the same candidates as building every operator, at about a twentieth of the cost on the test cases (supporting information, Section S10). Three conventions complete the procedure. Only operators of positive determinant are returned: $\mathbf{P}$ and $-\mathbf{P}$ predict the same metric, cell parameters carry no hand, and the convention is that both bases are right-handed, and an inverted geometry must be caught before the cell reaches this step. A fitted scale that changes the edge lengths by more than 10% is rejected, and so are two primitive cells whose volumes differ by more than that allows; a wrong centring symbol or a doubled axis lands here. The limit is far looser than any shape band, so isotropic non-isomorphism is never penalised. An obverse–reverse relation between two rhombohedral lattices keeps the volume and the hexagonal cell, so no cell-based test sees it; it is a matter for the intensities. A centred reference is first taken to a primitive cell (a rhombohedral cell on rhombohedral axes is already primitive) and the centring matrix composed with the operator found, which is then rational in the conventional basis. When the reference's widened set lies in the first layer the candidates can come from $\mathcal{U}$ directly (Table 7, step 3); $\mathbf{W}$ is then a column operator (only 768 of the 6096 transposes lie in $\mathcal{U}$), and since Selling reduction may return a left-handed basis, a proper $\mathbf{P}$ needs the 3048 matrices with $\det \mathbf{W} = \det B_R \det B_N$.

4.3. Choosing the band

The band is a tolerance on how much two measurements of one crystal may differ. For macromolecular crystals that difference is dominated by non-isomorphism (Crick & Magdoff, 1956; Juers & Matthews, 2001), not by the measurement error that the standard uncertainties of a cell describe. We measure non-isomorphism in the PDB (Section 4.4) and take the scale of the band from its upper percentiles, which depend on how a crystal form is delimited and on the Bravais type (Section 5.3). A candidate's price is read against that distribution: a class priced at the median non-isomorphism is as close to the reference as a typical second crystal of the same form, and the metric cannot exclude it. The tables of this paper use a price band of 7.4%, the implementation's default, the strain of one cell angle changed by 3° (Table 5). The value is illustrative: the deposited cells that set its scale were refined with symmetry imposed, which understates the non-isomorphism of cells measured without it (Section 6). As a price band it admits a pseudo-twofold only up to an obliquity of about 1.5°, so it is narrower than the lists of alternatives that data-processing programs test with the intensities: Pointless (Evans, 2006) and DIALS (Gildea & Winter, 2018) search lattice symmetry within 2° by default and the twin analysis of Phenix (Zwart *et al.*, 2005) within 3°, prices of 9.9 and 14.8%. The guarantee holds for whatever band is chosen; the band says which measurements it covers. Against non-isomorphism, 7.4% covers 95.7% of the pairs of one crystal form in the PDB, 9.9% covers 97.5% and 14.8% covers 99.4% (Section 5.3); the last two also reproduce the alternatives that 2° and 3° lattice-symmetry searches test, with the reference widened at 10.4 or 16.0% strain. For small-molecule data, whose cells are measured precisely and rarely non-isomorphous, a band of a few times the strain implied by the esds is the natural choice, and the implementation converts the esds of a cell to that strain; it is not tested here.

4.4. Tests

Constructed cases. The guarantee is tested on cases built where the closure predicts trouble (supporting information, Section S2). The benchmark has 780 primitive cells with a random metric error (median strain 0.1%) and a random proper reindexing: 75 measured in another shape of the same lattice, 660 near a specialisation, 20 whose only reduced superbase is reachable from the reference's only through a non-member, 16 needles and plates and 9 of the boundary sorts. The stress family puts independent errors on both cells of 60 references, scaled so that the true operator's price is 0.5 and 0.9 times the band. A further family has 300 lattices with two right angles sharing a vector or three right angles, each measured twice with an error of 0.1 to 2%, the reference given as its Niggli cell or in a random basis, and 60 lattices measured in a member that needs an entry of 3. In every case the true operator is known, and the test is whether it is among the candidates within the band. The same cases were also given to re-implementations of the searches that data-processing programs use (Section 5.1; supporting information, Section S8, Table S14).

The PDB. The snapshot of 206 214 cells deposited in the Protein Data Bank (Berman *et al.*, 2000; Zenodo record 22986222) is used twice. Non-isomorphism is measured between entries of one crystal form, taken as entries with the same space group whose identifiers are adjacent and share their first three characters, so that most pairs come from one deposition series; both cells are in the standard setting of the space group, so the identity is taken as the operator, and pairs above 20% are counted as different forms. The choices behind this proxy are varied as controls: the cut, identical cells, the weight of large deposition blocks, the Bravais type, and the cheapest operator in place of the identity. Pseudo-symmetry is measured by running the cctbx lattice-symmetry analysis at 3° on the primitive Niggli cell of every entry and pricing every proper operator of the highest lattice group that is neither a symmetry of the reported space group nor of the reported lattice's own holohedry. A null model shuffles the cell parameters within a space group. Deposited cells are refined in the reported space group, so they satisfy its symmetry exactly; what the analysis measures is how far real crystals lie from a higher lattice.

5. Results

Table 7
Reindexing a new cell against a reference within a price band ϵ.

step	action
1	Choose the price band ϵ (Section 4.3: 7.4% as an illustrative default against the non-isomorphism of deposited cells, more for unconstrained cells; 9.9% or 14.8% to list the alternatives that a 2° or 3° lattice-symmetry search would test). Reference: take the primitive cell, Selling-reduce it (B_R, F_0; a flip fires only at a conorm below -10^{-9} times the largest squared length, so exact zeros do not oscillate), and enumerate its widened set at strain $\eta = e^{\epsilon} - 1$ from the seven balls of Lemma A6; store the members, the operators $\mathbf{W}_k$ to them, and the smallest layer holding the set. A larger band needs the reference widened again.
2	New cell: take its primitive cell, with its own centring symbol (a rhombohedral cell on hexagonal axes reduced through its rhombohedral primitive cell, as for the reference), and Selling-reduce it once (B_N). Refuse if the primitive volumes differ by more than the scale limit allows, between 0.9^3 and 1.1^3 for a 10% limit on the edges.
3	Candidates: $\mathbf{P} = B_R\mathbf{W}_k\mathbf{Q}B_N^{-1}$ for every member and relabelling $\mathbf{Q}$ (equation 5); if the set lies in the first layer, equivalently $\mathbf{P} = B_R\mathbf{W}B_N^{-1}$ for the 3048 $\mathbf{W} \in \mathcal{U}$ (column operators) with $\det\mathbf{W} = \det B_R \det B_N$.
4	Keep $\det\mathbf{P} > 0$, price each by the geodesic distance after the best common scale, and keep those within ϵ whose scale changes the edges by at most 10% (supporting information, Section S10).
5	Group the candidates by the left action of the proper Laue group (point group for anomalous data); if the space group is itself in question, group by the proper holohedry of the reference's Bravais lattice, or report the candidates ungrouped. Report each class with its cheapest operator and price, read the prices against the non-isomorphism of the crystal's Bravais type (Table 10), and flag the result when the cheapest class costs more than its 95th percentile.
6	When more than one class lies within the band, decide with the intensities (Brehm & Diederichs, 2014; Gildea & Winter, 2018).

5.1. Completeness on measured cells

The worked cell. The measured cell 30.02, 39.97, 50.03 Å, 90.25°, 92.05°, 89.80° is constructed for this test: it has the edges of the 100° cell of Section 3.2, but β is set to 92°, so that the lattice lies near an orthorhombic one, and the cell is then perturbed by a metric error of strain 0.82%. Its widened closure grows with the strain: one member at 0.1%, two at 0.3%, three at 0.6%, the six of the ring from 1%, and sixteen at 7.4% (Fig. 2; the members, as flip words with their conorms and depths, are in Table S6). The six members of the ring enter below 0.8%; the other ten members of the closure of the nearby orthorhombic lattice enter between 3.8 and 5.1%.

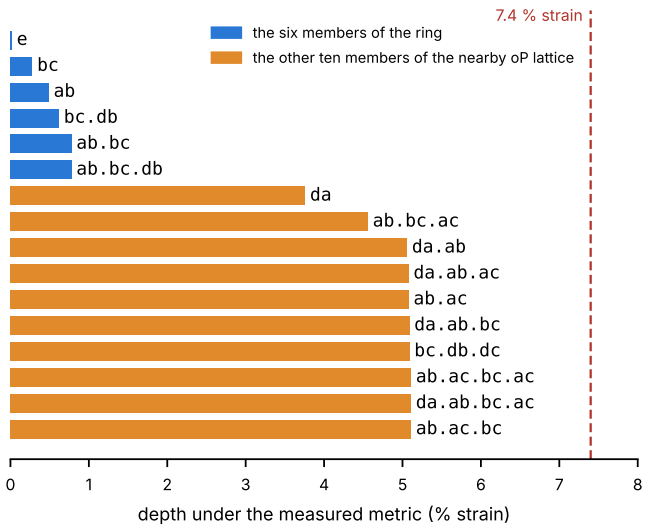


Figure 2
The sixteen members of the widened closure of the measured worked cell by depth: the six of the ring enter below 0.8% strain, the other ten members of the nearby orthorhombic lattice's closure between 3.8 and 5.1%.

Needles and plates. The exact enumeration and brute force agree on a needle 5.58 × 5.21 × 131.97 Å, $\gamma = 91.12°$, at strains from 0.05 to 5.2%, on two plates at 5.2%, where the widened sets have 20 and 28 members, and on 240 random needles, plates, isotropic and monoclinic cells. For very thin cells at large bands the ball of Lemma A6 grows quickly, and the implementation refuses balls above 5000 vectors with a message.

Where the 6096 matrices suffice. Table 8 gives the radius r_1 for six cells, and r_2, the same radius for the second layer $\mathcal{U}\mathcal{U}F_0$ of 2755 superbases. For the worked cell and a 10 × 10 × 30 Å needle no superbase outside the first layer comes within 25%, so the lookup of $\mathcal{U}$ is complete at every band in use. For needles and plates r_1 falls with the ratio of the shortest to the longest edge, to 4–6% for the thinnest cells, and the second layer roughly doubles it; there the lookup is complete only for bands below r_1, and the exact enumeration is needed beyond. For the BLDConograph sample cells and a lysozyme cell (supporting information, Table S5 and Section S6) r_1 reaches the cap, and so it does for 99.4% of the 697 benchmark references and 98.6% of 2000 PDB cells drawn at random; every one of those 2000 has r_1 above the 7.7% that certifies the default band. The archive has a thin tail: 3HR2 and 3HQV (40 × 27 × 678 Å, $r_1 = 6.0\%$) and 6BZP (13.8 × 4.9 × 101.9 Å, 6.8%) lie below 7.7%, where the enumeration uses the second layer, and a placeholder cell of 1 × 1 × 70.8 Å (1CGM) is refused by the cap on the ball. Widening a reference and reindexing a new cell take a few hundredths of a second for ordinary cells, and about a minute for a 300 × 300 × 3 Å plate (timings in the supporting information, Section S7).

Table 8
Radii in strain (%): r_1, the smallest strain at which a superbase outside the first layer $\mathcal{U}F_0$ becomes obtuse, and r_2 for the second layer $\mathcal{U}\mathcal{U}F_0$, searched up to 25%. Each value is certified in exact rational arithmetic, to the digits shown, for the reduced conorms of the cell as stored in double precision, by a dual bound for every superbase outside the layer within the cap and a feasible point for the nearest one (supporting information, Section S7).

cell	r_1	r_2
worked cell, 30.02 × 39.97 × 50.03 Å, 90.25, 92.05, 89.80°	> 25	> 25
needle 10 × 10 × 30 Å	> 25	> 25
needle 10 × 10 × 100 Å	13.9	> 25
needle 5.58 × 5.21 × 131.97 Å, $\gamma = 91.12°$	5.6	11.1
plate 100 × 100.1 × 5 Å, 90.03, 89.98, 90.05°	5.0	14.1
plate 150 × 120 × 4.8 Å, $\beta = 90.4°$	4.3	10.0

The guarantee on the test cases. The closure found the true operator in every case of every family, as Lemma A7 guarantees: 780 of 780 on the benchmark, 120 of 120 in the stress cases, 600 of 600 near two- and three-zero faces and 60 of 60 in members that need an entry of 3. The lookup of $\mathcal{U}$ found it in every case too, including 12 benchmark cases on three needle and plate references whose r_1 lies below the 7.7% that would certify it; that is not assured: against the $100 \times 100.1 \times 5$ Å plate of Table 8 ($r_1 = 5\%$), a new cell whose true operator costs 6.1% is found by the enumeration and missed by the lookup. *Existing searches.* Re-implementations of the searches in use were run on the same cases (supporting information, Table S14). The $\{-1, 0, 1\}$ window between the two Niggli cells and the cctbx similarity search found the true operator in every case of every family; the lattice-symmetry pipeline missed 1 of 780 benchmark cases, 10 of 120 stress cases and 8 of 600 near two- and three-zero faces. A $\{-1, 0, 1\}$ window between the two Selling-reduced cells misses operators where Table 4 says it can (supporting information, Section S8).

5.2. Indexing classes

Table 9 gives the classes for six references against the same crystal measured again with a 0.1% error in the length scale, which the scale fit absorbs, and a relabelling. For the exactly symmetric references the counts are the classical indexing ambiguities of merohedral crystals (Yeates, 1997; Dauter, 2003; Evans, 2006, 2011): one class for $P4_32_12$, four for $P3$, two for $P3_121$ and $P2_13$. The price adds the pseudo-merohedral classes: a $40.0 \times 40.3 \times 80$ Å $P2_12_12_1$ cell and the worked cell in $P1$ each have a second class, at 2.11% and 1.55%, twice, to first order, the cells' distances of 1.06% and 0.78% to the sorts the twofolds would make exact, from the twofold that would make the cell tetragonal or monoclinic. With proper operators only, anomalous data split no class of a chiral crystal; for an acentric crystal that is not chiral the point group replaces the Laue group, and $P4mm$ goes from one class to two and $P31m$ from two to four.

Table 9
Distinct indexings of six references at a band of 7.4% (the reference widened at 7.7%): the order of the band's group of proper self-maps H_ϵ and of the proper rotations $\mathcal{L}^+$ of the Laue group, the number of classes among the candidates, and the price of each class. The new cell is the reference with a 0.1% scale error, which the scale fit absorbs, relabelled.

reference	space group	$\lvert H_\epsilon \rvert$	$\lvert \mathcal{L}^+ \rvert$	classes	prices (%)
79.1 × 79.1 × 37.9 Å	$P4_32_12$	8	8	1	0
$a = b = 50$, $c = 80$ Å	$P3$	12	3	4	0, 0, 0, 0
$a = b = 50$, $c = 80$ Å	$P3_121$	12	6	2	0, 0
$a = 60$ Å	$P2_13$	24	12	2	0, 0
40.0 × 40.3 × 80 Å	$P2_12_12_1$	8	4	2	0, 2.11
worked cell	$P1$	2	1	2	0, 1.55

What the metric does not decide. The cheapest candidate lies in the true operator's coset of the exact lattice symmetries in 742 of the 780 benchmark cases, and in each of these it is the true operator itself. The 38 that fail are the 20 difference-wall cases and 18 of the 660 cases near a specialisation; in each the cheapest candidate is the true operator composed with a pseudo-symmetry, of the nearby orthorhombic corner for the difference-wall cases, that costs less than the true operator at that error. In the stress cases the cheapest candidate is in the true coset in 68 of 120, and is the true operator itself in 40. When a pseudo-symmetry costs less than the error, the price orders the candidates but cannot pick the true one; how often that happens in real crystals is the subject of the next section.

5.3. Application: the price of pseudo-symmetry in the PDB

Non-isomorphism. Between deposited cells of one crystal form (45 222 pairs) the price after the scale fit has a median of 0.98%, a 90th percentile of 4.2% and a 95th of 6.7% (95% block-bootstrap intervals over deposition blocks, 4.1–4.3 and 6.5–7.0%; Efron, 1979; Künsch, 1989); without the scale fit the median is 1.6%. The upper percentiles depend on how a crystal form is delimited: with the cut between forms at 5, 10, 20 and 50% the 90th percentile is 2.5, 3.4, 4.2 and 5.6% and the 95th 3.1, 4.7, 6.7 and 11.9%. Dropping the 1 464 pairs of identical cells changes little (4.3 and 6.9%), and weighting each deposition block equally, so that fragment-screening series of up to 35 pairs count once, gives 5.0 and 8.0%. By reported Bravais type (Table 10) the 90th and 95th percentiles are 2.6 and 4.0% for tP, 2.9 and 4.6% for hP, 4.7 and 7.4% for oP, 5.1 and 8.4% for mC and 5.4 and 9.0% for mP; for the cubic types they are 0, the scale fit removing the whole difference. For aP they are 7.7 and 11.7% with the identity as the operator, and 6.1 and 9.9% with the cheapest operator, since two triclinic deposits of one crystal are often in different reduced settings; for mP, mC and oP the cheapest operator changes them by 0.5 points or less (supporting information, Section S9). The scale a reindexing band has to cover for macromolecular data is therefore a few per cent: at the 20% cut, 2.6–6.1% at the 90th percentile and 4–10% at the 95th across the non-cubic types (aP with the cheapest operator); the standard uncertainties of a single cell describe its measurement error alone and do not reach it. The default band of 7.4% lies above the pooled 95th percentile and above the 90th percentile of every type except aP with the identity as the operator (7.7%; 95% interval 6.4–9.5%). The intervals of every percentile are in

the supporting information (Table S15); the oF row rests on 31 pairs in 22 blocks and is not informative, and the cubic rows are zero by construction.

Table 10

Non-isomorphism in the PDB by reported Bravais type: the price (%, after the scale fit) between deposited cells of one crystal form, taken as id-adjacent entries of one space group below 20% unfitted strain, with the identity as the operator; for the low-symmetry types also the 90th and 95th percentiles with the cheapest operator. The cells are deposited cells refined with the reported symmetry imposed, so only their free parameters carry error: cells from indexing, or from serial and multi-crystal data, need wider bands. For the cubic types the scale fit removes the whole difference; the oF row rests on 22 deposition blocks. Bootstrap intervals are in Table S15.

type	pairs	median	90th	95th	90th / 95th, cheapest operator
aP	902	2.06	7.66	11.68	6.14 / 9.89
mP	7 077	1.27	5.38	9.02	5.29 / 8.80
mC	4 932	1.49	5.06	8.40	4.82 / 7.90
oP	13 615	1.11	4.70	7.43	4.65 / 7.39
oS	2 271	1.01	3.57	5.50	—
oI	1 353	1.22	5.12	7.91	—
oF	31	1.99	5.87	7.54	—
tP	5 174	0.52	2.59	4.00	—
tI	805	0.65	2.76	4.91	—
hR	1 270	0.68	3.54	5.86	—
hP	6 817	0.57	2.87	4.56	—
cP	363	0.00	0.00	0.00	—
cI	385	0.00	0.00	0.00	—
cF	211	0.00	0.00	0.00	—

Pseudo-symmetry. Pseudo-symmetry is common in macromolecular crystals and a frequent source of errors in space-group assignment (Zwart *et al.*, 2008). Le Page's analysis at 3° finds a higher lattice than the reported one for 43 201 of the 206 141 entries analysed (21%); 31 placeholder cells with an edge below 2 Å and 42 cells the cctbx analysis did not accept are left out. For 545 (0.3%) the higher lattice is exact. Operators of the reported lattice's own holohedry, taken from its Bravais type, which a merohedral space group such as $R3$ or $P4_1$ leaves out of its point group, are not counted as pseudo-symmetry. The operators are those of the cctbx analysis, priced as in this paper. For the 42 656 entries without an exact higher lattice the largest price matches $2\sqrt{2}\,\mathrm{artanh}(\sin\delta)$ for the group's largest obliquity δ (ratio 1.000 at the 10th, 50th and 90th percentiles), consistent with Proposition A9 when a twofold is the most expensive operator, and a check of the code, and the cheapest operator has a median price of 6.7%. Fig. 3 sets these prices against non-isomorphism: the cheapest alternative lies within the 90th percentile of non-isomorphism for 35% of the pseudo-symmetric entries and within the 95th for 50%. Over all entries, 7.5% (95% interval 7.4–7.7%) have an alternative, exact or priced, within 4.2% and 10.6% (10.4–11.0%) within 6.7%; with the band of the entry's own Bravais type, the pooled band standing in for a type with fewer than 200 deposition blocks, 8.6% (8.4–9.0%) and 12.5% (12.0–13.0%).

Against chance. Shuffling the free cell parameters among the entries of one space group, with and without keeping cells of one volume together, shows that pseudo-symmetry within 3° is reached by chance almost as often as in real cells, while near-exact pseudo-symmetry is not chance: the share whose cheapest alternative costs at most 1% is 10% for real $P1$ cells against 2–3% shuffled, and 9% against 5–6% for $P2_1$ (supporting information, Section S9). The obliquities of the priced alternatives are spread nearly evenly below 3°, so median prices of 6–7.5% follow from the 3° window, and the match of the largest price to $2\sqrt{2}\,\mathrm{artanh}(\sin\delta)$ is a check of the implementation against Proposition A9 rather than a finding.

What the cell decides. An alternative within the band is one the cell does not exclude; it is not one the cell prefers. Reindexing the second cell of 3 000 randomly chosen same-form pairs against the first, with the first cell's space group and a band of 7.4%, a wrong class lies within the band at a positive price in 7.9% of pairs, and at exactly the identity's price in 17%: the classical merohedral ambiguities within the reported lattice type, such as $P3$ or $P4$ in a lattice of higher holohedry, which no cell can resolve. A wrong class is strictly cheaper than the identity's in 36 pairs (1.2%; 0.8–1.6% with deposition blocks resampled): in 12 by at most 0.05 points, below the precision of a deposited cell; in 20 through a relabelling of axes, which mixes deposits that use different axis orders, where the identity is not the true operator, with genuine near-ties of pseudo-tetragonal cells; and in 4 otherwise (0.13%, 0.03–0.27%). Telling these apart needs the deposited models, which we did not compare. In 125 pairs (4.2%) the two cells differ by more than 7.4%, measurements outside what that band covers: the upper tail of non-isomorphism (4.3% of all 45 222 same-form pairs, median price 10.7%), which a wider band covers (Section 4.3); in 9 of them (0.3%) only a wrong class lies within 7.4%. So for about one pair in thirteen the cell leaves a priced alternative open, and it ranks a wrong indexing first in between about 0.1% and 1.2% of pairs, depending on how the relabellings are read. This matters when cells are compared, in reindexing against a reference or in merging, not for a single data set.

Where. Monoclinic and triclinic entries carry most of the share: 18% of the 35 364 mP entries and 16% of the 8 491 aP entries have an alternative within 4.2%, against 7% of oP and 1% of tP entries. Among the transitions, mP to hP is nearly always undecidable from the cell (median price 0.71%, 91% within 4.2%), the configuration in which pseudo-merohedral twinning is possible (Parsons, 2003); for most transitions the median price of the cheapest alternative is 6–7.5%, so the sort decides how often a pseudo-symmetry occurs, not what it costs (supporting information, Tables S8 and S9).

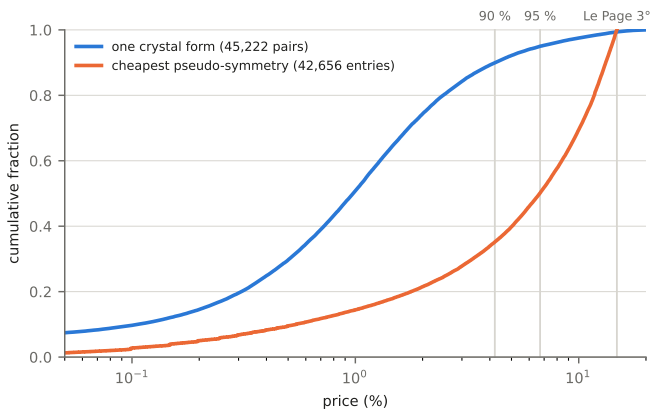


Figure 3
Non-isomorphism against pseudo-symmetry in the PDB: cumulative distributions of the price between deposited cells of one crystal form (blue; one value per pair) and of the price of the cheapest pseudo-symmetric alternative of each entry for which Le Page's 3° finds a higher lattice that is not exact (orange; one value per entry). The orange curve ends near 14.8% by construction, since an obliquity of at most 3° prices a twofold at most so. Vertical lines: the 90th and 95th percentiles of non-isomorphism, 4.2 and 6.7%, and the price of a twofold at Le Page's 3°. Prices are on a logarithmic axis, so pairs of price zero (identical cells, and most cubic pairs) lie off it; the per-type percentiles are in Table 10.

5.4. Large entries, and why the Niggli window works

When does reduction put a zero on d? The entries are read in the basis of the first cell, so the question is about the reference's reduced superbase, and near a lattice with three right angles the Niggli type answers it (Proposition A11). A type I measurement, whose Niggli cell has three acute angles, has no obtuse orthogonal cell, so near enough to the right angles every reduction ends on a face-diagonal cell whose zeros form a path through d, whatever basis it starts from: 402 of 402 reductions of measured near-oP cells from random bases, and 67 of 67 from the Niggli cell (supporting information, Section S7). A type II Niggli cell is itself obtuse with its small conorms off d, so reduction started from it stays there (83 of 83), but reduction started from another basis put a near-zero conorm on d in 370 of 498 cases. Near a lattice with two right angles a type II Niggli cell kept its right angles on ab, ac, bc in 519 of 519 measured near-mP cells, as Proposition A11(i) requires; from a type I Niggli cell the outcome depends on the order of flips: reduction put at least one near-zero conorm on d in 614 of 681, and both near-zero conorms in 316. Which type a measurement gets is decided by the signs of its small errors; in that sample 57% of near-mP measurements were type I.

Why the Niggli window works. The true operator between the two Niggli cells was computed directly, with no search, for 100 lattices on each of the 38 faces and on each of the 24 sorts, measured twice with independent errors of total strain 0.5 to 7.7% and the new cell in a random basis: its entries were in $\{-1, 0, 1\}$ in all 24 800 cases, and in 3 040 further cases on strongly anisotropic lattices, the square root of the metric's eigenvalue ratio up to 98 (supporting information, Section S9). Between the two Selling-reduced cells the same operators needed an entry of 2 or 3 in 3 765 cases, all on lattices with two or three right angles, and none once composed with the lattice's point group. A Niggli cell is built on short vectors and has no fourth vector d, so it does not meet the zeros on d that make the Selling entries large. Proposition A12 proves this for exact lattices: two Niggli cells of metrics with one zero pattern, written relative to one obtuse superbase, differ by a matrix with entries in $\{-1, 0, 1\}$. It proves it also for two measurements within a radius ρ of one lattice, which ranges from 1.9% for a generic triclinic lattice to 50% for a cubic one (supporting information, Table S21); the errors above reach 7.7%, beyond ρ for some lattices, so the observation reaches further than the proof. Corollary A13 answers the converse question, how to get from a reduced cell to the Niggli cell: from every obtuse superbase, in each of its 24 ordered bases, the Niggli cell is a $\{-1, 0, 1\}$ matrix away, and the Niggli vectors are three of its seven partial sums, up to sign. Whichever member of the closure a reduction lands on, the Niggli cell can be read off it without a second reduction, and no member has to be avoided. The guarantee of this paper depends on neither result.

6. Discussion

Completeness. Reindexing as described here returns exactly the proper operators whose price lies within the band. The operators between reduced cells of any lattice lie in one finite set, every element of which is needed for some lattice; for a measured cell the reduced cells within a tolerance are finitely many and are listed; so the true operator is among the candidates whenever its price is within the band, and no method that uses the two cells alone can return fewer and keep that guarantee. The candidates are grouped into classes that differ by pseudo-symmetries whose price is known, and the classes, not the operators, are the distinct indexing hypotheses. The searches that programs use, from unit-cell similarity searches to lattice-symmetry pipelines, find the true operator in ordinary cases; on the cases tested here the $\{-1, 0, 1\}$ window between Niggli cells and the cctbx similarity search found it in every case, and our re-implementation of the lattice-symmetry pipeline missed 19 of 1560 (Section 5.1). The theory says why a fixed list can work: below r_1 every candidate is one of the 6096 matrices of $\mathcal{U}$, so any fixed list that contains their images has the same reach there, and between Niggli cells those images lie within $\{-1, 0, 1\}$, as Proposition A12 proves for exact lattices and for measurements within a stated radius (Section 5.4); beyond r_1 the enumeration supplies what a fixed list cannot promise.

Structure, and bounds. The closure of a lattice is a connected groupoid and so is determined, up to equivalence, by the point group; what is new is that all its matrices, for every lattice, are computed once, as $\mathcal{U}$, the neighbours of Voronoi's principal domain (Voronoi, 1908; Schürmann, 2009). A measured cell has no such structure of its own, only near-symmetries that need not compose within the tolerance. Their prices are the values of the displacement function of the action of $GL_3(\mathbb{Z})$ on metrics (Bridson & Haefliger, 1999, Chapter II.6), finite in number below any bound because the action is properly discontinuous, and bounded on both sides by the symmetries of the nearest symmetric lattice (Proposition A10(v)). Thin cells lie far out

in the cusps of the space of lattices, where a shear that adds a short vector to a long one moves the metric very little; the number of near-symmetries there has no bound independent of the cell, which is why the ball of Lemma A6 grows for needles and plates at large bands.

Pseudo-symmetry and non-isomorphism. The price of a pseudo-symmetric alternative is a geodesic distance and can be compared directly with the distance between two crystals of one form. Where the alternative costs less than ordinary non-isomorphism, no metric criterion can exclude it and the intensities must decide, as in twinning and indexing-ambiguity tests (Yeates, 1997; Evans, 2006, 2011; Brehm & Diederichs, 2014; Kabsch, 2014; Liu & Spence, 2014; Gildea & Winter, 2018). The direct measure is the reindexing of pairs of one crystal form: for about one pair in thirteen the cell leaves a priced alternative open within 7.4% (Section 5.3). Counting entries instead, about one in twelve has a pseudo-symmetric alternative priced within the 90th percentile of non-isomorphism, an upper screen that does not compare two measurements. Both are measures of metric ambiguity under a proxy for one crystal form, adjacent identifiers of one space group, which clustering on the cells or the intensities (Giordano *et al.*, 2012; Foadi *et al.*, 2013; Diederichs, 2017) would sharpen; neither counts data that were actually misindexed, which would need the intensities of the deposited data. At about 5% of price per degree of Le Page obliquity, the 90th percentile of non-isomorphism corresponds to about 0.85°, and Le Page's 3°, a tolerance for assigning a Bravais lattice, prices a twofold at two and a half to six times it; for twofolds the price depends on the obliquity and not on the sort, and the sort decides how often each kind of pseudo-symmetry occurs (Table S8).

The Niggli cell. The instability of the Niggli cell (Gruber, 1973, 1989; Andrews *et al.*, 1980; Andrews & Bernstein, 1988, 2014; Mighell, 2001) is the same phenomenon in part. At a zero scalar product the Niggli cell changes between type I and type II, and the closure handles that boundary as it handles any zero conorm; at a tie such as $a = b$, and at Niggli's special conditions, the Niggli basis changes while the Selling-reduced superbase does not, being only relabelled (Andrews *et al.*, 2019*a*). Proposition A11 adds the sign condition under which a type I Niggli cell leads to a different Selling-reduced superbase. Tolerant Niggli reduction (Křivý & Gruber, 1976; Grosse-Kunstleve *et al.*, 2004) returns one of these cells, not the set.

Related work. In the S^6 picture of Andrews and Bernstein, the reduced presentations of a lattice are already related by boundary transformations, and CS6Dist and NCDist measure distances between lattices; they return a distance, not an operator. Kurlin's root invariant identifies a lattice up to isometry and varies continuously with it. This paper adds the explicit universal set of matrices between reduced presentations, the exact bound on their entries by Voronoi type, completeness under a tolerance certified for each cell by r_1, and a price that does not depend on the basis. Oishi-Tomiyasu (2012, 2024) gives a provably sufficient list of basis transforms for Bravais-lattice determination under an error bound; $\mathcal{U}$ is the corresponding list for reindexing between reduced cells, complete for every exact lattice and, below r_1, for every measured one. Measuring approximate symmetry by the distance to the symmetric configurations, as Proposition A9 does for lattices, is the idea of continuous symmetry measures for molecules (Zabrodsky *et al.*, 1992). The same computations give the distance of a cell to each of the 24 sorts, compared with BLDConograph and SELLA in the supporting information (Section S6).

Scope. "Complete" in this paper means: every proper unimodular operator whose price is at most the band is among the candidates, for a band the user supplies. This holds for primitive relations of determinant ± 1, below a strain of 25%, for every cell whose ball of Lemma A6 holds at most 5000 lattice vectors (thinner cells at large bands are refused with a message), with membership in the widened set decided at one consistent relative tolerance of 10^{-9} (Section 3.4); the 6096 matrices alone give the same set when the widened set lies in the first layer. Sub- and superlattice relations (a doubled axis, determinant ± 2 or more) are outside the closure and are refused with a message; an obverse–reverse relation between two rhombohedral lattices keeps the cell and is not seen at all; the choice within a class is left to the intensities. The default band of 7.4% is illustrative. The PDB measurements that set its scale describe deposited cells, refined with their reported symmetry imposed, whose free parameters alone carry error; cells from indexing before symmetry is imposed, or from serial and multi-crystal data, carry error in all six parameters and need wider bands, which a calibration on unconstrained per-crystal cells would give. The guarantee holds for whatever band is chosen.

7. Conclusions

Every integer matrix between two reduced cells of any lattice lies in one set of 6096 matrices, and which of them occur depends only on which conorms vanish; both statements and the converse are checked in Lean. For a measured cell the reduced cells within a strain tolerance are finitely many and are listed, so reindexing returns exactly the proper operators within a stated price band, the true one among them whenever its price is within the band; prices place the candidates between cosets of the near-symmetries, bounded by the symmetries of the nearest symmetric lattice. In practice (Table 7): widen each reference once at the chosen band, Selling-reduce each new data set once, group the candidates by the Laue group, and read each class's price against the non-isomorphism of its Bravais type (Table 10). The band decides which measurements are covered: for deposited cells of one crystal form, 7.4% covers 95.7% of the pairs and 14.8%, which also reproduces the alternatives of a 3° lattice-symmetry search, 99.4%, and unconstrained cells need more. Within the band the intensities decide.

Data availability. The methods are implemented in the open-source package agentsg (Zwart, 2026; https://github.com/phzwart/agentsg), version 0.5.0. It reduces cells, widens a reference under a band, returns the candidate operators between a reference and a new cell with their prices, groups them into indexing classes, provides the operator set

$\mathcal{U}$, and reads the Niggli cell off the Selling-reduced cell. The 6096 operators of $\mathcal{U}$ are also deposited as a self-describing JSON file that states what the operators are, their convention, their determinants and how to apply them between two Selling-reduced cells. The PDB cell snapshot is on Zenodo (doi:10.5281/zenodo.22986222). An archived release of the repository, with a DOI, will accompany publication [DOI to be inserted]; it contains the machine-checked proofs and the computational checks of every computable statement and every number in this paper, which the supporting information lists by section, table and figure (Section S7, Table S13).

[To be completed by the author: funding, e.g. the U.S. Department of Energy, Office of Science, programme and award, under Contract No. DE-AC02-05CH11231 to Lawrence Berkeley National Laboratory; acknowledgements. The use of AI is disclosed in the footnote to Section 1.]

Appendix A
The results and how they are proved

This appendix states the results the paper rests on, says what each does in the argument, and outlines its proof. The statements in full, with complete proofs under the same numbers, are in the supporting information (Section S12); Table 6 records which are machine-checked and which rest on an exhaustive computation, and every one has a computational check (supporting information, Section S7). The status words, here and in Table 6, mean the following. *Proved*: a complete argument written out by hand (supporting information, Section S12), using only the cited results and the earlier statements. *Machine-checked*: the statement, as formalised, is a theorem of the proof assistant Lean 4 with Mathlib, verified by its kernel; where the proof contains a finite computation, that computation is run by compiled code inside Lean, which adds the Lean compiler to the trusted base, and the axioms each theorem uses are listed in the supporting information (Table S12). A machine check covers the formal statement, which may be narrower than the prose: the window $[-7,7]$ for Lemma A2, for instance, or any isometric action for Propositions A9(i) and A10. *Computer-assisted*: the proof reduces the statement to finitely many cases, and a program decides every case in exact integer or rational arithmetic. A case is set aside only with a certificate that is itself checked exactly (for Proposition A12, a Farkas certificate that a system of linear inequalities has no solution), and floating point may only enlarge the set of cases examined, never shrink it. *Certified*: a quantity bracketed in exact arithmetic, by a feasible point from above and a dual bound from below. *Computed numerically*: obtained in floating point, without a certificate, such as ρ in Proposition A12. *First order*: exact in the limit of small strain, with an error of higher order. *Observed*: seen in every case of a test, without a proof. Every statement, whatever its status, also has a computational check: an exhaustive enumeration, a comparison with brute force, or seeded random samples (supporting information, Section S7). A sampled check can refute a statement but does not prove it.

Vectors are written in the basis (a,b,c) of a superbase (d,a,b,c), $d=-a-b-c$; a *partial sum* is the sum of some of the four vectors, and up to sign the seven partial sums lie one in each parity class other than 2Λ.

Lemma A1 (class excess identity). For any superbase, the squared length of a partial sum minus that of another vector of its parity class is $-4\sum p_{ij}\gamma_{ij}$, where the γ_{ij} are non-negative integers fixed by the two vectors' coefficients. *Role.* It turns "shortest in its class" into the vanishing of integers on the slots with non-zero conorm, so which vectors are shortest, and hence which superbases are obtuse, depends only on the zero pattern. *Proof.* Selling's identity $|x|^2 = \sum_{i<j} p_{ij}(c_i - c_j)^2$ for $x=\sum c_i v_i$.

Lemma A2 (window). The vectors of any obtuse superbase of a lattice have coefficients in $[-6,7]$ in the basis of any other. *Role.* It reduces the completeness of $\mathcal{U}$ to a finite enumeration. *Proof.* Lemma A1 applied to the partial sums of both superbases, which are all shortest in their classes, along the slots with non-zero conorm, which connect the four vectors. Machine-checked.

The enumeration. For each of the 38 zero patterns, list the vectors of the window that are shortest in their classes for every metric with that pattern (Lemma A1), and form the superbases all of whose partial sums are on the list: these are exactly the obtuse superbases of every lattice with the pattern. In their 48 labellings they give the $6096 = 48 \times 127$ matrices of $\mathcal{U}$, and 1, 2, 3, 6 and 16 superbases for no zero, one zero, two zeros on disjoint pairs, two zeros sharing a vector and three zeros, independently of Kurlin's (2022) lists. The 127 superbases of the first layer are also counted by hand, pattern by pattern, as $1+6+36+12+72$.

Lemma A3 (generation). Every matrix of a closure, and every symmetry, is a word of boundary transformations (Selling steps at zero conorms) followed by a relabelling. *Role.* The closure is reached by the boundary transformations of Andrews *et al.* (2023), symmetries included; the sorts (A5) rest on it. *Proof.* Perturb the zero conorms so that the target superbase becomes the only obtuse one, and run Selling's reduction from the other: every step it takes is at a conorm that vanishes for the original metric.

Proposition A4 (the strain bounds every dot product). If a metric error $\Delta\mathbf{G}$ has strain at most η, then $|u^{\mathsf{T}}\Delta\mathbf{G}\,v| \le \eta|u||v|$ for all u, v. *Role.* A metric error becomes a band on every conorm; Lemma A6, Proposition A12 and the necessary condition of the reindexing workflow use it. *Proof.* The spectral norm is at most the Frobenius norm.

Definition and enumeration A5 (sorts as closed subspaces). For each zero pattern, the sorts are the generic points of the subspaces of conorm space that are fixed by a set of the closure's matrices, each a set of equalities among conorms; enumerated, they are Delone's 24 sorts. *Role.* The table of sorts follows from the closure (Section 2.3). *Proof.* The fixed subspaces and the symmetry sets form a Galois connection, whose closed subspaces are finitely many intersections, enumerated pattern by pattern.

Lemma A6 (the widened set is finite). Every vector of a super-

base within strain $\eta < 1$ of obtuse lies in a ball of squared radius $\frac{1+\eta}{1-\eta}$ times the shortest squared length of its class. *Role.* The widened set is enumerated exactly, with no window of coordinates assumed. *Proof.* Proposition A4 with $u = v$, and the minimality of the partial sums of an obtuse superbase.

Lemma A7 (widening one side suffices). An operator of price at most ϵ carries the new cell's reduced superbase into the reference's widened set at $\eta = e^{\epsilon} - 1$. *Role.* This is the guarantee: widening the reference alone finds every candidate within the band. *Proof.* A geodesic distance ϵ bounds the strain by $e^{\epsilon} - 1$, and the operator is an isometry of both.

Remark A8 (layers, and the lookup of $\mathcal{U}$). The layers $\mathcal{N}_1 = \mathcal{U}F_0, \mathcal{N}_k = \mathcal{U}\mathcal{N}_{k-1}$ and the adjacency radius r_1: below r_1, every candidate is $B_R \mathbf{W} B_N^{-1}$ for one of the 6096 column operators $\mathbf{W}$ of $\mathcal{U}$ (3048 for proper operators). *Role.* It says when a fixed list of matrices suffices. *Proof.* Lemma A7, the completeness of $\mathcal{U}$ and its closure under relabelling. The radii of Table 8 are certified in exact arithmetic; no geometric lower bound on r_1 is proved.

Proposition A9 (the price of a symmetry). A twofold's price is twice the distance to the metrics it fixes, and $2\sqrt{2}\,\mathrm{artanh}(\sin\delta)$ for a twofold of obliquity δ; for orders 3, 4 and 6 the ratio is $\sqrt{3}$, $\{\sqrt{2}, 2\}$ and $\{1, \sqrt{3}\}$, to first order. *Role.* It ties the price to Le Page's obliquity (3° is 14.8%) and to the distances to the sorts. *Proof.* The action of g is an isometry; for a twofold the geodesic midpoint is fixed (Bhatia, 2007); the obliquity formula is an explicit eigenvalue computation, and the first-order ratios are rotation angles of $X \mapsto g^{\mathsf{T}} X g$.

Proposition A10 (candidates and cosets). Prices are subadditive under products of self-maps; the candidates lie between two cosets of the reference's cheap self-maps; a gap in the prices makes those self-maps a group and the candidates one coset; and the near-symmetries of a reference lie between the point group of a nearby symmetric lattice and its self-maps within a slightly larger band. *Role.* It is why the candidates group into indexing classes (Section 5.2). *Proof.* The triangle inequality for the geodesic distance, on which unimodular matrices act by isometries.

Proposition A11 (Niggli type and the zeros on d). A type II Niggli cell (a, b, c) gives an obtuse superbase (d, a, b, c) whose small conorms are off d; a type I Niggli cell gives none of the orthogonal superbases, so near three right angles a type I measurement reduces, from any basis, to a cell whose zeros form a path through d. *Role.* It explains when reduction puts a zero on d and so when entries of 2 and 3 appear (Section 5.4). *Proof.* Niggli's conditions bound the conorms of d from below, and the sign of $(a \cdot b)(a \cdot c)(b \cdot c)$ does not change when the signs of basis vectors change. Parts (i) and (ii) machine-checked.

Proposition A12 (the Niggli window; computer-assisted). (i) Two Niggli cells of metrics with one zero pattern, written relative to one obtuse superbase, differ by a matrix with entries in $\{-1, 0, 1\}$. (ii) So do the Niggli cells of two measurements within a strain ρ of one lattice; ρ is set by the gap between the shortest and the next vectors of the parity classes and is 1.9% for a generic triclinic lattice, 10.4% for the lattice of Section 3.2 and 50% for a cubic one (supporting information, Table S21). *Role.* It explains why the $\{-1, 0, 1\}$ search between Niggli cells succeeds (Section 5.4). *Proof.* The vectors of a Niggli cell are shortest in their parity classes, so a Niggli cell is a basis of the at most 26 vectors that Lemma A1 lists for the pattern, satisfying Minkowski's conditions, which are linear in the conorms. Over all 38 patterns the 8 256 bases that satisfy them somewhere are found by linear programming, the other 166 656 are excluded by exact certificates, and all 3 849 216 pairs are checked. ρ is computed numerically.

Corollary A13 (from a reduced cell to the Niggli cell; computer-assisted). For every lattice and every obtuse superbase, the Niggli vectors are three of its seven partial sums, up to sign. Equivalently, from each of its 24 ordered bases the Niggli cell is reached by one of 1 344 fixed matrices (672 up to sign), those whose columns are among the 14 signed partial sums. *Role.* Every member of the closure is a landing spot for the Niggli cell, which is read off the reduced cell without a second reduction. *Proof.* The finite check of Proposition A12, extended to every obtuse superbase of every pattern and each of its ordered bases (6 856 704 matrices). The list is complete; that it is minimal is not shown.

References

Andrews, L. C. & Bernstein, H. J. (1988). *Acta Cryst.* A**44**, 1009–1018.
Andrews, L. C. & Bernstein, H. J. (2014). *J. Appl. Cryst.* **47**, 346–359.
Andrews, L. C. & Bernstein, H. J. (2023). *Acta Cryst.* A**79**, 499–503.
Andrews, L. C., Bernstein, H. J. & Pelletier, G. A. (1980). *Acta Cryst.* A**36**, 248–252.
Andrews, L. C., Bernstein, H. J. & Sauter, N. K. (2019a). *Acta Cryst.* A**75**, 115–120.
Andrews, L. C., Bernstein, H. J. & Sauter, N. K. (2019b). *Acta Cryst.* A**75**, 593–599.
Andrews, L. C., Bernstein, H. J. & Sauter, N. K. (2023). SELLA – a program for determining Bravais lattice types. arXiv:2303.03122.
Berman, H. M., Westbrook, J., Feng, Z., Gilliland, G., Bhat, T. N., Weissig, H., Shindyalov, I. N. & Bourne, P. E. (2000). *Nucleic Acids Res.* **28**, 235–242.
Bhatia, R. (2007). *Positive Definite Matrices.* Princeton: Princeton University Press.
Brandt, H. (1927). *Math. Ann.* **96**, 360–366.
Brehm, W. & Diederichs, K. (2014). *Acta Cryst.* D**70**, 101–109.
Bridson, M. R. & Haefliger, A. (1999). *Metric Spaces of Non-Positive Curvature.* Berlin: Springer.
Burzlaff, H. & Zimmermann, H. (1985). *Z. Kristallogr.* **170**, 247–262.
Burzlaff, H. & Zimmermann, H. (2006). *International Tables for Crystallography*, Vol. A, 5th ed., edited by Th. Hahn, ch. 9.1, pp. 742–749. Dordrecht: Springer.
Cliffe, M. J. & Goodwin, A. L. (2012). *J. Appl. Cryst.* **45**, 1321–1329.
Conway, J. H. & Sloane, N. J. A. (1992). *Proc. R. Soc. London Ser. A*, **436**, 55–68.
Crick, F. H. C. & Magdoff, B. S. (1956). *Acta Cryst.* **9**, 901–908.
Dauter, Z. (2003). *Acta Cryst.* D**59**, 2004–2016.
de Moura, L. & Ullrich, S. (2021). *Automated Deduction – CADE 28*, Lecture Notes in Computer Science, Vol. 12699, pp. 625–635. Cham: Springer.
de Wolff, P. M. & Gruber, B. (1991). *Acta Cryst.* A**47**, 29–36.
Delaunay, B. (1933). *Z. Kristallogr.* **84**, 109–149.
Delone, B. N., Galiulin, R. V. & Shtogrin, M. I. (1975). *J. Sov. Math.* **4**, 79–156.
Diederichs, K. (2017). *Acta Cryst.* D**73**, 286–293.
Efron, B. (1979). *Ann. Stat.* **7**, 1–26.

Engel, P., Michel, L. & Senechal, M. (2004). *Lattice Geometry.* Preprint. Bures-sur-Yvette: Institut des Hautes Études Scientifiques.

Evans, P. R. (2006). *Acta Cryst.* D**62**, 72–82.

Evans, P. R. (2011). *Acta Cryst.* D**67**, 282–292.

Flack, H. D. (1987). *Acta Cryst.* A**43**, 564–568.

Foadi, J., Aller, P., Alguel, Y., Cameron, A., Axford, D., Owen, R. L., Armour, W., Waterman, D. G., Iwata, S. & Evans, G. (2013). *Acta Cryst.* D**69**, 1617–1632.

Gildea, R. J. & Winter, G. (2018). *Acta Cryst.* D**74**, 405–410.

Giordano, R., Leal, R. M. F., Bourenkov, G. P., McSweeney, S. & Popov, A. N. (2012). *Acta Cryst.* D**68**, 649–658.

Grosse-Kunstleve, R. W., Sauter, N. K. & Adams, P. D. (2004). *Acta Cryst.* A**60**, 1–6.

Grosse-Kunstleve, R. W., Sauter, N. K., Moriarty, N. W. & Adams, P. D. (2002). *J. Appl. Cryst.* **35**, 126–136.

Gruber, B. (1973). *Acta Cryst.* A**29**, 433–440.

Gruber, B. (1989). *Acta Cryst.* A**45**, 123–131.

Hahn, Th. (2006). Editor. *International Tables for Crystallography*, Vol. A, 5th ed. Dordrecht: Springer.

Hattne, J., Echols, N., Tran, R., Kern, J., Gildea, R. J., Brewster, A. S., Alonso-Mori, R., Glöckner, C., Hellmich, J., Laksmono, H., Sierra, R. G., Lassalle-Kaiser, B., Lampe, A., Han, G., Gul, S., DiFiore, D., Milathianaki, D., Fry, A. R., Miahnahri, A., White, W. E., Schafer, D. W., Seibert, M. M., Koglin, J. E., Sokaras, D., Weng, T.-C., Sellberg, J., Latimer, M. J., Glatzel, P., Zwart, P. H., Grosse-Kunstleve, R. W., Bogan, M. J., Messerschmidt, M., Williams, G. J., Boutet, S., Messinger, J., Zouni, A., Yano, J., Bergmann, U., Yachandra, V. K., Adams, P. D. & Sauter, N. K. (2014). *Nat. Methods*, **11**, 545–548.

Higgins, P. J. (1971). *Notes on Categories and Groupoids.* London: Van Nostrand Reinhold.

Himes, V. L. & Mighell, A. D. (1987). *Acta Cryst.* A**43**, 375–384.

Juers, D. H. & Matthews, B. W. (2001). *J. Mol. Biol.* **311**, 851–862.

Kabsch, W. (2010). *Acta Cryst.* D**66**, 125–132.

Kabsch, W. (2014). *Acta Cryst.* D**70**, 2204–2216.

Křivý, I. & Gruber, B. (1976). *Acta Cryst.* A**32**, 297–298.

Künsch, H. R. (1989). *Ann. Stat.* **17**, 1217–1241.

Kurlin, V. (2022). *A complete isometry classification of 3-dimensional lattices.* arXiv:2201.10543.

Le Page, Y. (1982). *J. Appl. Cryst.* **15**, 255–259.

Liebschner, D., Afonine, P. V., Baker, M. L., Bunkóczi, G., Chen, V. B., Croll, T. I., Hintze, B., Hung, L.-W., Jain, S., McCoy, A. J., Moriarty, N. W., Oeffner, R. D., Poon, B. K., Prisant, M. G., Read, R. J., Richardson, J. S., Richardson, D. C., Sammito, M. D., Sobolev, O. V., Stockwell, D. H., Terwilliger, T. C., Urzhumtsev, A. G., Videau, L. L., Williams, C. J. & Adams, P. D. (2019). *Acta Cryst.* D**75**, 861–877.

Liu, H. & Spence, J. C. H. (2014). *IUCrJ*, **1**, 393–401.

mathlib Community (2020). *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs*, pp. 367–381. New York: ACM.

Mighell, A. D. (2001). *J. Res. Natl Inst. Stand. Technol.* **106**, 983–995.

Moakher, M. (2005). *SIAM J. Matrix Anal. Appl.* **26**, 735–747.

Neff, P., Eidel, B. & Martin, R. J. (2016). *Arch. Ration. Mech. Anal.* **222**, 507–572.

Niggli, P. (1928). *Kristallographische und strukturtheoretische Grundbegriffe.* Handbuch der Experimentalphysik, Vol. 7, Part 1. Leipzig: Akademische Verlagsgesellschaft.

Oishi-Tomiyasu, R. (2012). *Acta Cryst.* A**68**, 525–535.

Oishi-Tomiyasu, R. (2024). *Acta Cryst.* A**80**, 339–350.

Parsons, S. (2003). *Acta Cryst.* D**59**, 1995–2003.

Pennec, X., Fillard, P. & Ayache, N. (2006). *Int. J. Comput. Vis.* **66**, 41–66.

Santoro, A. & Mighell, A. D. (1970). *Acta Cryst.* A**26**, 124–127.

Sauter, N. K., Grosse-Kunstleve, R. W. & Adams, P. D. (2004). *J. Appl. Cryst.* **37**, 399–409.

Schlenker, J. L., Gibbs, G. V. & Boisen, M. B. (1978). *Acta Cryst.* A**34**, 52–54.

Schürmann, A. (2009). *Computational Geometry of Positive Definite Quadratic Forms.* University Lecture Series, Vol. 48. Providence: American Mathematical Society.

Selling, E. (1874). *J. Reine Angew. Math.* **77**, 143–229.

Voronoi, G. (1908). *J. Reine Angew. Math.* **134**, 198–287.

Weinstein, A. (1996). *Notices Am. Math. Soc.* **43**, 744–752.

White, T. A., Kirian, R. A., Martin, A. V., Aquila, A., Nass, K., Barty, A. & Chapman, H. N. (2012). *J. Appl. Cryst.* **45**, 335–341.

Winter, G., Waterman, D. G., Parkhurst, J. M., Brewster, A. S., Gildea, R. J., Gerstel, M., Fuentes-Montero, L., Vollmar, M., Michels-Clark, T., Young, I. D., Sauter, N. K. & Evans, G. (2018). *Acta Cryst.* D**74**, 85–97.

Yeates, T. O. (1997). *Methods Enzymol.* **276**, 344–358.

Zabrodsky, H., Peleg, S. & Avnir, D. (1992). *J. Am. Chem. Soc.* **114**, 7843–7851.

Zwart, P. H. (2026). agentsg: a space-group and lattice engine, version 0.5.0. https://github.com/phzwart/agentsg.

Zwart, P. H., Grosse-Kunstleve, R. W. & Adams, P. D. (2005). *CCP4 Newsl. Protein Crystallogr.* **43**, contribution 10.

Zwart, P. H., Grosse-Kunstleve, R. W., Lebedev, A. A., Murshudov, G. N. & Adams, P. D. (2008). *Acta Cryst.* D**64**, 99–107.